

#Predict the price of the Uber ride from a given pickup point to the agreed drop-off location. Perform following tasks:

1. Pre-process the dataset.
2. Identify outliers.
3. Check the correlation.
4. Implement linear regression and random forest regression models.
5. Evaluate the models and compare their respective scores like R2, RMSE, etc. Dataset link: <https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>

#Importing the required libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

#importing the dataset

```
df = pd.read_csv("uber.csv")
```

1. Pre-process the dataset.

```
df.head()
```

| | Unnamed: 0 | | key | fare_amount |
|---|-------------------|-------------------------|------|---------------------|
| | pickup_datetime \ | | | |
| 0 | 24238194 | 2015-05-07 19:52:06 UTC | 7.5 | 2015-05-07 19:52:06 |
| 1 | 27835199 | 2009-07-17 20:04:56 UTC | 7.7 | 2009-07-17 20:04:56 |
| 2 | 44984355 | 2009-08-24 21:45:00 UTC | 12.9 | 2009-08-24 21:45:00 |
| 3 | 25894730 | 2009-06-26 08:22:21 UTC | 5.3 | 2009-06-26 08:22:21 |
| 4 | 17610152 | 2014-08-28 17:47:00 UTC | 16.0 | 2014-08-28 17:47:00 |

| | pickup_longitude | pickup_latitude | dropoff_longitude |
|---|--------------------|-----------------|-------------------|
| | dropoff_latitude \ | | |
| 0 | -73.999817 | 40.738354 | -73.999512 |
| | 40.723217 | | |
| 1 | -73.994355 | 40.728225 | -73.994710 |
| | 40.750325 | | |
| 2 | -74.005043 | 40.740770 | -73.962565 |
| | 40.772647 | | |
| 3 | -73.976124 | 40.790844 | -73.965316 |
| | 40.803349 | | |
| 4 | -73.925023 | 40.744085 | -73.973082 |
| | 40.761247 | | |

```

    passenger_count
0                1
1                1
2                1
3                3
4                5

```

`df.info()` *#To get the required information of the dataset*

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null  int64
1   key                   200000 non-null  object
2   fare_amount           200000 non-null  float64
3   pickup_datetime       200000 non-null  object
4   pickup_longitude      200000 non-null  float64
5   pickup_latitude       200000 non-null  float64
6   dropoff_longitude     199999 non-null  float64
7   dropoff_latitude      199999 non-null  float64
8   passenger_count       200000 non-null  int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB

```

`df.columns` *#To get number of columns in the dataset*

```

Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
      'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
      'dropoff_latitude', 'passenger_count'],
      dtype='object')

```

`df = df.drop(['Unnamed: 0', 'key'], axis= 1)` *#To drop unnamed column as it isn't required*

`df.head()`

```

    fare_amount      pickup_datetime  pickup_longitude
pickup_latitude \
0      7.5  2015-05-07 19:52:06 UTC      -73.999817
40.738354
1      7.7  2009-07-17 20:04:56 UTC      -73.994355
40.728225
2     12.9  2009-08-24 21:45:00 UTC      -74.005043
40.740770
3      5.3  2009-06-26 08:22:21 UTC      -73.976124
40.790844
4     16.0  2014-08-28 17:47:00 UTC      -73.925023
40.744085

```

```

    dropoff_longitude  dropoff_latitude  passenger_count

```

| | | | |
|---|------------|-----------|---|
| 0 | -73.999512 | 40.723217 | 1 |
| 1 | -73.994710 | 40.750325 | 1 |
| 2 | -73.962565 | 40.772647 | 1 |
| 3 | -73.965316 | 40.803349 | 3 |
| 4 | -73.973082 | 40.761247 | 5 |

df.shape *#To get the total (Rows,Columns)*

(200000, 7)

df.dtypes *#To get the type of each column*

```
fare_amount      float64
pickup_datetime  object
pickup_longitude float64
pickup_latitude  float64
dropoff_longitude float64
dropoff_latitude float64
passenger_count  int64
dtype: object
```

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 200000 entries, 0 to 199999

Data columns (total 7 columns):

| # | Column | Non-Null Count | Dtype |
|---|-------------------|-----------------|---------|
| 0 | fare_amount | 200000 non-null | float64 |
| 1 | pickup_datetime | 200000 non-null | object |
| 2 | pickup_longitude | 200000 non-null | float64 |
| 3 | pickup_latitude | 200000 non-null | float64 |
| 4 | dropoff_longitude | 199999 non-null | float64 |
| 5 | dropoff_latitude | 199999 non-null | float64 |
| 6 | passenger_count | 200000 non-null | int64 |

dtypes: float64(5), int64(1), object(1)

memory usage: 10.7+ MB

df.describe() *#To get statistics of each columns*

| | fare_amount | pickup_longitude | pickup_latitude | |
|-------|---------------|------------------|-----------------|---|
| count | 200000.000000 | 200000.000000 | 200000.000000 | |
| mean | 11.359955 | -72.527638 | 39.935885 | - |
| std | 9.901776 | 11.437787 | 7.720539 | |
| min | -52.000000 | -1340.648410 | -74.015515 | - |
| 25% | 6.000000 | -73.992065 | 40.734796 | - |
| 75% | 11.359955 | -72.527638 | 39.935885 | |

| | | | | |
|-------------|------------|------------|-------------|---|
| 50% | 8.500000 | -73.981823 | 40.752592 | - |
| 73.980093 | | | | |
| 75% | 12.500000 | -73.967153 | 40.767158 | - |
| 73.963659 | | | | |
| max | 499.000000 | 57.418457 | 1644.421482 | |
| 1153.572603 | | | | |

| | dropoff_latitude | passenger_count |
|-------|------------------|-----------------|
| count | 19999.000000 | 20000.000000 |
| mean | 39.923890 | 1.684535 |
| std | 6.794829 | 1.385997 |
| min | -881.985513 | 0.000000 |
| 25% | 40.733823 | 1.000000 |
| 50% | 40.753042 | 1.000000 |
| 75% | 40.768001 | 2.000000 |
| max | 872.697628 | 208.000000 |

Filling Missing values

```
df.isnull().sum()
```

```
fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude 1
dropoff_latitude 1
passenger_count  0
dtype: int64
```

```
df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(),inplace = True)
```

```
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(),inplace = True)
```

```
df.isnull().sum()
```

```
fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude 0
dropoff_latitude 0
passenger_count  0
dtype: int64
```

```
df.dtypes
```

```
fare_amount      float64
pickup_datetime  object
pickup_longitude  float64
pickup_latitude  float64
```

```

dropoff_longitude    float64
dropoff_latitude     float64
passenger_count      int64
dtype: object

```

Column pickup_datetime is in wrong format (Object). Convert it to DateTime Format

```

df.pickup_datetime = pd.to_datetime(df.pickup_datetime,
errors='coerce')

```

```

df.dtypes

```

```

fare_amount          float64
pickup_datetime      datetime64[ns, UTC]
pickup_longitude     float64
pickup_latitude      float64
dropoff_longitude    float64
dropoff_latitude     float64
passenger_count      int64
dtype: object

```

To segregate each time of date and time

```

df= df.assign(hour = df.pickup_datetime.dt.hour,
              day= df.pickup_datetime.dt.day,
              month = df.pickup_datetime.dt.month,
              year = df.pickup_datetime.dt.year,
              dayofweek = df.pickup_datetime.dt.dayofweek)

```

```

df.head()

```

```

   fare_amount  pickup_datetime  pickup_longitude
pickup_latitude \
0      7.5 2015-05-07 19:52:06+00:00      -73.999817
40.738354
1      7.7 2009-07-17 20:04:56+00:00      -73.994355
40.728225
2     12.9 2009-08-24 21:45:00+00:00      -74.005043
40.740770
3      5.3 2009-06-26 08:22:21+00:00      -73.976124
40.790844
4     16.0 2014-08-28 17:47:00+00:00      -73.925023
40.744085

```

```

   dropoff_longitude  dropoff_latitude  passenger_count  hour  day
month \
0      -73.999512      40.723217          1      19      7
5
1      -73.994710      40.750325          1      20     17
7
2      -73.962565      40.772647          1      21     24
8
3      -73.965316      40.803349          3       8     26

```

```

6
4      -73.973082      40.761247      5      17      28
8

```

```

    year  dayofweek
0  2015          3
1  2009          4
2  2009          0
3  2009          4
4  2014          3

```

```

# drop the column 'pickup_datetime' using drop()
# 'axis = 1' drops the specified column

```

```

df = df.drop('pickup_datetime',axis=1)

```

```

df.head()

```

```

    fare_amount  pickup_longitude  pickup_latitude
dropoff_longitude \
0          7.5      -73.999817      40.738354      -73.999512
1          7.7      -73.994355      40.728225      -73.994710
2         12.9      -74.005043      40.740770      -73.962565
3          5.3      -73.976124      40.790844      -73.965316
4         16.0      -73.925023      40.744085      -73.973082

```

```

    dropoff_latitude  passenger_count  hour  day  month  year
dayofweek
0      40.723217          1      19      7      5  2015
3
1      40.750325          1      20     17      7  2009
4
2      40.772647          1      21     24      8  2009
0
3      40.803349          3       8     26      6  2009
4
4      40.761247          5     17     28      8  2014
3

```

```

df.dtypes

```

```

fare_amount      float64
pickup_longitude  float64
pickup_latitude   float64
dropoff_longitude  float64

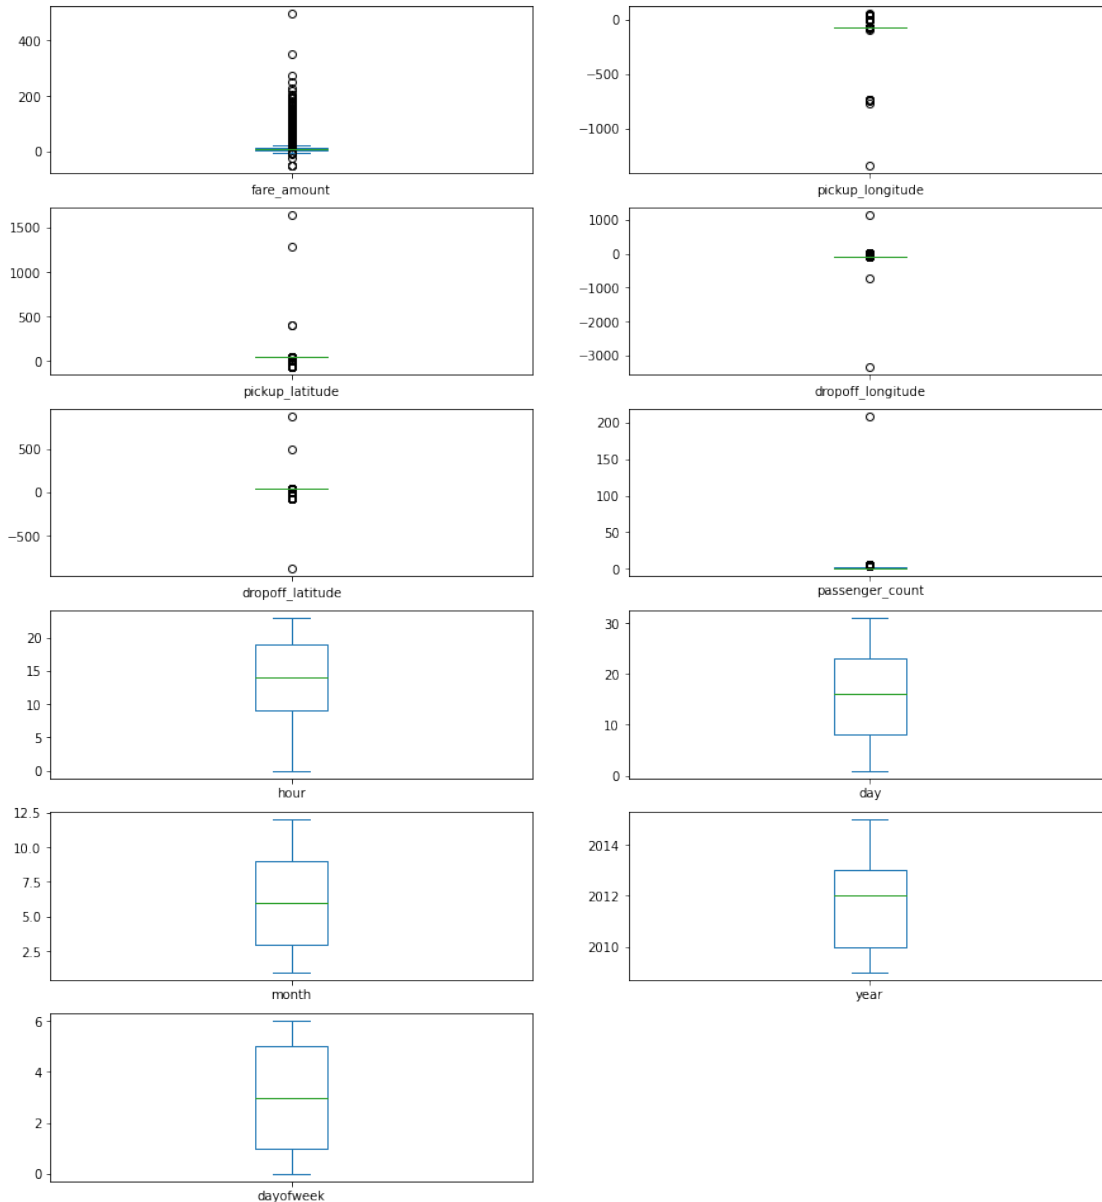
```

```
dropoff_latitude    float64
passenger_count     int64
hour                int64
day                 int64
month               int64
year                int64
dayofweek           int64
dtype: object
```

Checking outliers and filling them

```
df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20))
#Boxplot to check the outliers
```

```
fare_amount          AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude     AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude      AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude     AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude     AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count      AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour                 AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day                  AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month                AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year                 AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek            AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object
```



#Using the InterQuartile Range to fill the values

```
def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1-1.5*IQR
    upper_whisker = Q3+1.5*IQR
    df[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
    return df1

def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
    return df1
```



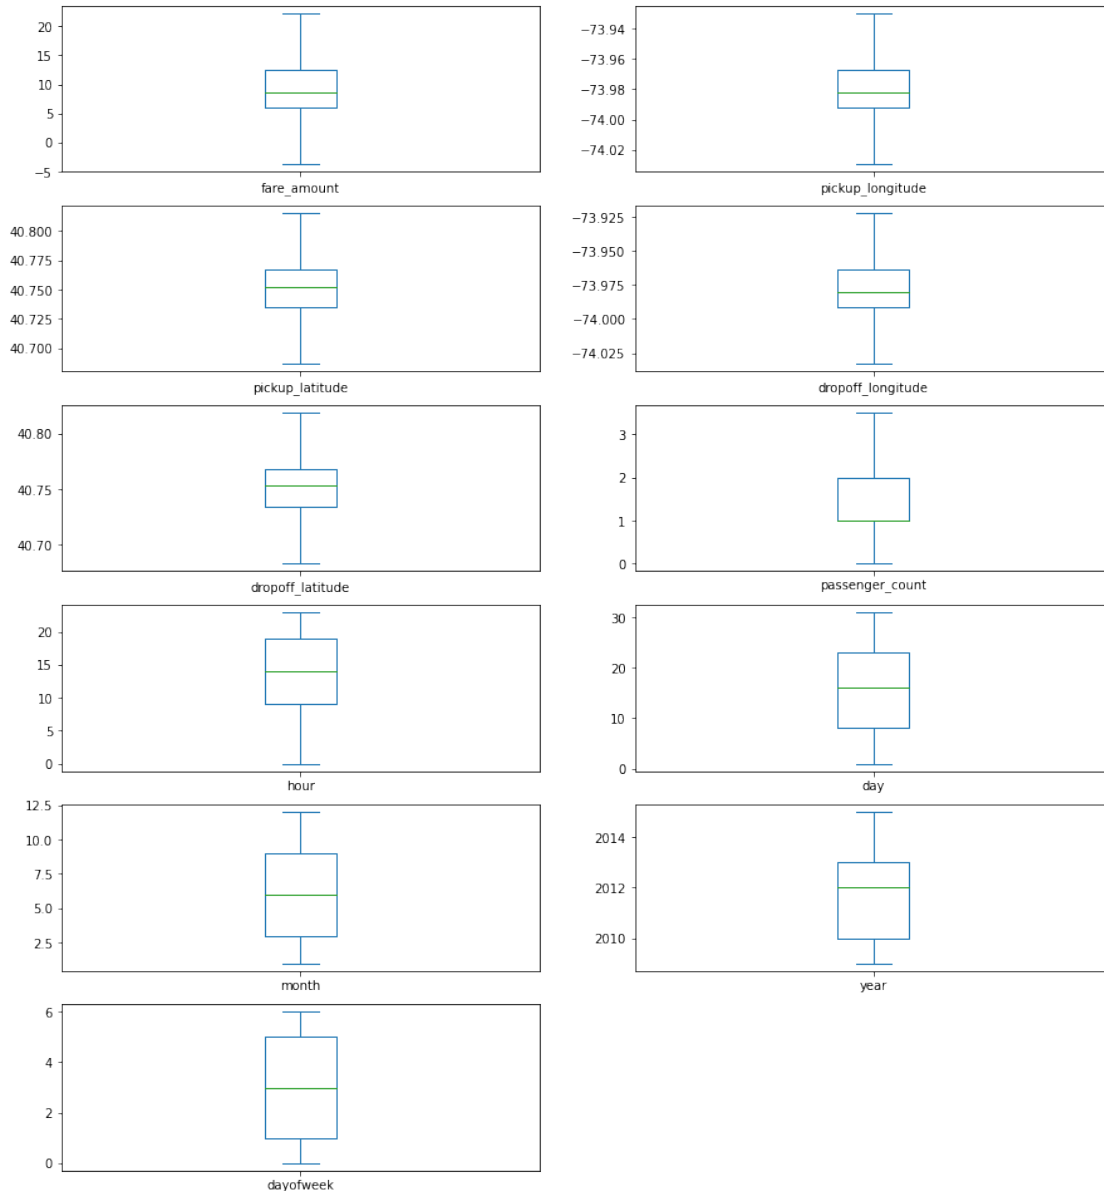
```

df = treat_outliers_all(df , df.iloc[: , 0::])

df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20))
#Boxplot shows that dataset is free from outliers

fare_amount          AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude     AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude      AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude    AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude     AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count      AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour                 AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day                  AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month                AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year                 AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek            AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object

```



```
#pip install haversine
import haversine as hs #Calculate the distance using Haversine to
calculate the distance between to points. Can't use Eucladian as it is
for flat surface.
travel_dist = []
for pos in range(len(df['pickup_longitude'])):
    long1,lati1,long2,lati2 = [df['pickup_longitude']
[pos],df['pickup_latitude'][pos],df['dropoff_longitude']
[pos],df['dropoff_latitude'][pos]]
    loc1=(lati1,long1)
    loc2=(lati2,long2)
    c = hs.haversine(loc1,loc2)
    travel_dist.append(c)
```

```
print(travel_dist)
df['dist_travel_km'] = travel_dist
df.head()
```

IOPub data rate exceeded.
The notebook server will temporarily stop sending output
to the client in order to avoid crashing it.
To change this limit, set the config variable
`--NotebookApp.iopub_data_rate_limit`.

Current values:
NotebookApp.iopub_data_rate_limit=1000000.0 (bytes/sec)
NotebookApp.rate_limit_window=3.0 (secs)

| | fare_amount | pickup_longitude | pickup_latitude | dropoff_longitude \ |
|---|-------------|------------------|-----------------|---------------------|
| 0 | 7.5 | -73.999817 | 40.738354 | -73.999512 |
| 1 | 7.7 | -73.994355 | 40.728225 | -73.994710 |
| 2 | 12.9 | -74.005043 | 40.740770 | -73.962565 |
| 3 | 5.3 | -73.976124 | 40.790844 | -73.965316 |
| 4 | 16.0 | -73.929786 | 40.744085 | -73.973082 |

| | dropoff_latitude | passenger_count | hour | day | month | year |
|---|------------------|-----------------|------|-----|-------|------|
| 0 | 40.723217 | 1.0 | 19 | 7 | 5 | 2015 |
| 3 | | | | | | |
| 1 | 40.750325 | 1.0 | 20 | 17 | 7 | 2009 |
| 4 | | | | | | |
| 2 | 40.772647 | 1.0 | 21 | 24 | 8 | 2009 |
| 0 | | | | | | |
| 3 | 40.803349 | 3.0 | 8 | 26 | 6 | 2009 |
| 4 | | | | | | |
| 4 | 40.761247 | 3.5 | 17 | 28 | 8 | 2014 |
| 3 | | | | | | |

| | dist_travel_km |
|---|----------------|
| 0 | 1.683325 |
| 1 | 2.457593 |
| 2 | 5.036384 |
| 3 | 1.661686 |
| 4 | 4.116088 |

```
#Uber doesn't travel over 130 kms so minimize the distance
df= df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
print("Remaining observastions in the dataset:", df.shape)
```

Remaining observastions in the dataset: (200000, 12)

```
#Finding inccorect latitude (Less than or greater than 90) and longitude (greater than or less than 180)
incorrect_coordinates = df.loc[(df.pickup_latitude > 90) |
(df.pickup_latitude < -90) |
(df.dropoff_latitude > 90) |
(df.dropoff_latitude < -90) |
(df.pickup_longitude > 180) |
(df.pickup_longitude < -180) |
(df.dropoff_longitude > 90) |
(df.dropoff_longitude < -90)
]
```

```
df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')
```

```
df.head()
```

| | fare_amount | pickup_longitude | pickup_latitude | dropoff_longitude \ |
|---|-------------|------------------|-----------------|---------------------|
| 0 | 7.5 | -73.999817 | 40.738354 | -73.999512 |
| 1 | 7.7 | -73.994355 | 40.728225 | -73.994710 |
| 2 | 12.9 | -74.005043 | 40.740770 | -73.962565 |
| 3 | 5.3 | -73.976124 | 40.790844 | -73.965316 |
| 4 | 16.0 | -73.929786 | 40.744085 | -73.973082 |

| | dropoff_latitude | passenger_count | hour | day | month | year |
|---|------------------|-----------------|------|-----|-------|------|
| 0 | 40.723217 | 1.0 | 19 | 7 | 5 | 2015 |
| 3 | | | | | | |
| 1 | 40.750325 | 1.0 | 20 | 17 | 7 | 2009 |
| 4 | | | | | | |
| 2 | 40.772647 | 1.0 | 21 | 24 | 8 | 2009 |
| 0 | | | | | | |
| 3 | 40.803349 | 3.0 | 8 | 26 | 6 | 2009 |
| 4 | | | | | | |
| 4 | 40.761247 | 3.5 | 17 | 28 | 8 | 2014 |
| 3 | | | | | | |

| | dist_travel_km |
|---|----------------|
| 0 | 1.683325 |
| 1 | 2.457593 |


```
2      5.036384
3      1.661686
4      4.116088
```

```
df.isnull().sum()
```

```
fare_amount      0
pickup_longitude 0
pickup_latitude  0
dropoff_longitude 0
dropoff_latitude 0
passenger_count  0
hour             0
day             0
month           0
year            0
dayofweek       0
dist_travel_km   0
dtype: int64
```

```
sns.heatmap(df.isnull()) #Free for null values
```

```
<AxesSubplot:>
```



```
corr = df.corr() #Function to find the correlation
```

corr

| | fare_amount | pickup_longitude | pickup_latitude | \ |
|-------------------|-------------|------------------|-----------------|---|
| fare_amount | 1.000000 | 0.154069 | -0.110842 | |
| pickup_longitude | 0.154069 | 1.000000 | 0.259497 | |
| pickup_latitude | -0.110842 | 0.259497 | 1.000000 | |
| dropoff_longitude | 0.218675 | 0.425619 | 0.048889 | |
| dropoff_latitude | -0.125898 | 0.073290 | 0.515714 | |
| passenger_count | 0.015778 | -0.013213 | -0.012889 | |
| hour | -0.023623 | 0.011579 | 0.029681 | |
| day | 0.004534 | -0.003204 | -0.001553 | |
| month | 0.030817 | 0.001169 | 0.001562 | |
| year | 0.141277 | 0.010198 | -0.014243 | |
| dayofweek | 0.013652 | -0.024652 | -0.042310 | |
| dist_travel_km | 0.786385 | 0.048446 | -0.073362 | |

| | dropoff_longitude | dropoff_latitude | |
|-------------------|-------------------|------------------|---|
| passenger_count \ | | | |
| fare_amount | 0.218675 | -0.125898 | |
| 0.015778 | | | |
| pickup_longitude | 0.425619 | 0.073290 | - |
| 0.013213 | | | |
| pickup_latitude | 0.048889 | 0.515714 | - |
| 0.012889 | | | |
| dropoff_longitude | 1.000000 | 0.245667 | - |
| 0.009303 | | | |
| dropoff_latitude | 0.245667 | 1.000000 | - |
| 0.006308 | | | |
| passenger_count | -0.009303 | -0.006308 | |
| 1.000000 | | | |
| hour | -0.046558 | 0.019783 | |
| 0.020274 | | | |
| day | -0.004007 | -0.003479 | |
| 0.002712 | | | |
| month | 0.002391 | -0.001193 | |
| 0.010351 | | | |
| year | 0.011346 | -0.009603 | - |
| 0.009749 | | | |
| dayofweek | -0.003336 | -0.031919 | |
| 0.048550 | | | |
| dist_travel_km | 0.155191 | -0.052701 | |
| 0.009884 | | | |

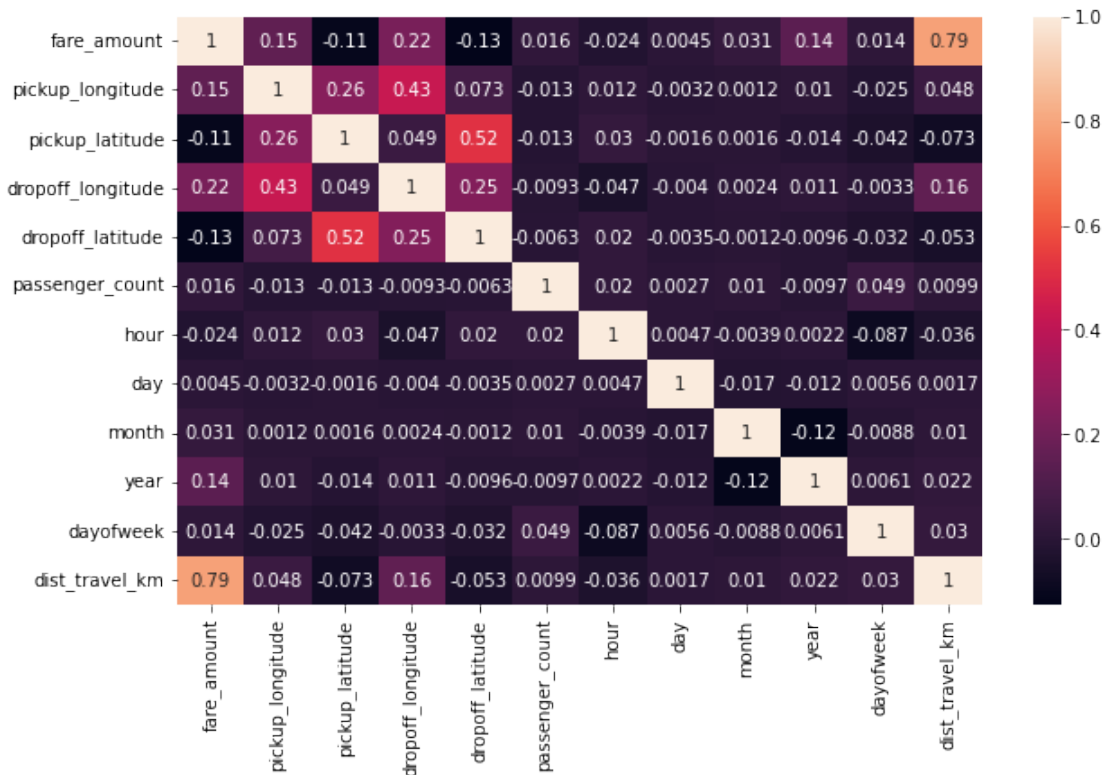
| | hour | day | month | year | |
|------------------|-----------|-----------|----------|-----------|-----------|
| dayofweek \ | | | | | |
| fare_amount | -0.023623 | 0.004534 | 0.030817 | 0.141277 | 0.013652 |
| pickup_longitude | 0.011579 | -0.003204 | 0.001169 | 0.010198 | -0.024652 |
| pickup_latitude | 0.029681 | -0.001553 | 0.001562 | -0.014243 | -0.042310 |

| | | | | | |
|-------------------|-----------|-----------|-----------|-----------|-----------|
| dropoff_longitude | -0.046558 | -0.004007 | 0.002391 | 0.011346 | -0.003336 |
| dropoff_latitude | 0.019783 | -0.003479 | -0.001193 | -0.009603 | -0.031919 |
| passenger_count | 0.020274 | 0.002712 | 0.010351 | -0.009749 | 0.048550 |
| hour | 1.000000 | 0.004677 | -0.003926 | 0.002156 | -0.086947 |
| day | 0.004677 | 1.000000 | -0.017360 | -0.012170 | 0.005617 |
| month | -0.003926 | -0.017360 | 1.000000 | -0.115859 | -0.008786 |
| year | 0.002156 | -0.012170 | -0.115859 | 1.000000 | 0.006113 |
| dayofweek | -0.086947 | 0.005617 | -0.008786 | 0.006113 | 1.000000 |
| dist_travel_km | -0.035708 | 0.001709 | 0.010050 | 0.022294 | 0.030382 |

| | dist_travel_km |
|-------------------|----------------|
| fare_amount | 0.786385 |
| pickup_longitude | 0.048446 |
| pickup_latitude | -0.073362 |
| dropoff_longitude | 0.155191 |
| dropoff_latitude | -0.052701 |
| passenger_count | 0.009884 |
| hour | -0.035708 |
| day | 0.001709 |
| month | 0.010050 |
| year | 0.022294 |
| dayofweek | 0.030382 |
| dist_travel_km | 1.000000 |

```
fig,axis = plt.subplots(figsize = (10,6))
sns.heatmap(df.corr(),annot = True) #Correlation Heatmap (Light values
means highly correlated)
```

```
<AxesSubplot:>
```



Dividing the dataset into feature and target values

```
x =
df[['pickup_longitude','pickup_latitude','dropoff_longitude','dropoff_latitude',
'passenger_count','hour','day','month','year','dayofweek','dist_travel_km']]
```

```
y = df['fare_amount']
```

Dividing the dataset into training and testing dataset

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size = 0.33)
```

Linear Regression

```
from sklearn.linear_model import LinearRegression
regression = LinearRegression()
```

```
regression.fit(X_train,y_train)
```

```
LinearRegression()
```

```
regression.intercept_ #To find the linear intercept
```

```
3683.734379131267
```

```
regression.coef_ #To find the linear coefficient
```

```
array([ 2.54690644e+01, -7.25031311e+00,  2.01910609e+01, -
1.81689799e+01,
```



```

        6.49318535e-02,  8.88740039e-03,  3.96976218e-03,
6.07701750e-02,
        3.64995448e-01, -3.34018868e-02,  1.84796864e+00])

prediction = regression.predict(X_test) #To predict the target values
print(prediction)

[ 6.92808422  5.50169187  7.29033891 ...  7.34427831 11.48600676
 8.04489363]

y_test
23033      8.0
166557     4.5
188533     8.0
175085     7.5
69692     11.4
...
22917      8.1
42396     12.9
25947      8.0
66067      8.5
20658      8.5
Name: fare_amount, Length: 66000, dtype: float64

```

Metrics Evaluation using R2, Mean Squared Error, Root Mean Squared Error

```

from sklearn.metrics import r2_score

r2_score(y_test,prediction)

0.6640797581905353

from sklearn.metrics import mean_squared_error

MSE = mean_squared_error(y_test,prediction)

MSE

9.92519776977491

RMSE = np.sqrt(MSE)

RMSE

3.15042818832217

```

Random Forest Regression

```

from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor(n_estimators=100) #Here n_estimators means
number of trees you want to build before making the prediction

rf.fit(X_train,y_train)

```

```
y_pred = rf.predict(X_test)
```

```
y_pred
```

Metrics evaluatin for Random Forest

```
R2_Random = r2_score(y_test,y_pred)
```

```
R2_Random
```

```
MSE_Random = mean_squared_error(y_test,y_pred)
```

```
MSE_Random
```

```
RMSE_Random = np.sqrt(MSE_Random)
```

```
RMSE_Random
```

Assignment 2

1. Classify the email using the binary classification method. Email Spam detection has two states: a) Normal State – Not Spam, b) Abnormal State – Spam. Use K-Nearest Neighbors and Support Vector Machine for classification. Analyze their performance. Dataset link: The emails.csv dataset on the Kaggle

<https://www.kaggle.com/datasets/balaka18/email-spam-classification-dataset-csv>

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn import metrics
```

```
df=pd.read_csv('emails.csv')
```

```
df.head()
```

```
      Email No.  the  to  ect  and  for  of   a  you  hou  ...  connevey
jay \
0      Email 1    0   0   1   0   0   0   2   0   0  ...         0
0
1      Email 2    8  13  24   6   6   2  102   1  27  ...         0
0
2      Email 3    0   0   1   0   0   0   8   0   0  ...         0
0
3      Email 4    0   5  22   0   5   1  51   2  10  ...         0
0
4      Email 5    7   6  17   1   5   2  57   0   9  ...         0
0
```

```
      valued  lay  infrastructure  military  allowing  ff  dry
Prediction
0          0   0                0          0          0   0   0
0
1          0   0                0          0          0   1   0
0
2          0   0                0          0          0   0   0
0
3          0   0                0          0          0   0   0
0
4          0   0                0          0          0   1   0
0
```

```
[5 rows x 3002 columns]
```

```

df.columns

Index(['Email No.', 'the', 'to', 'ect', 'and', 'for', 'of', 'a',
      'you', 'hou',
      ...,
      'connevey', 'jay', 'valued', 'lay', 'infrastructure',
      'military',
      'allowing', 'ff', 'dry', 'Prediction'],
      dtype='object', length=3002)

df.isnull().sum()

Email No.      0
the            0
to            0
ect           0
and           0
...
military      0
allowing      0
ff           0
dry          0
Prediction    0
Length: 3002, dtype: int64

df.dropna(inplace = True)

df.drop(['Email No.'],axis=1,inplace=True)
X = df.drop(['Prediction'],axis = 1)
y = df['Prediction']

from sklearn.preprocessing import scale
X = scale(X)
# split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 42)

##KNN classifier

from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=7)

knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

print("Prediction",y_pred)

Prediction [0 0 1 ... 1 1 1]

print("KNN accuracy = ",metrics.accuracy_score(y_test,y_pred))

KNN accuracy = 0.8009020618556701

```

```
print("Confusion matrix",metrics.confusion_matrix(y_test,y_pred))
```

```
Confusion matrix [[804 293]
 [ 16 439]]
```

SVM classifier

```
# cost C = 1
```

```
model = SVC(C = 1)
```

```
# fit
```

```
model.fit(X_train, y_train)
```

```
# predict
```

```
y_pred = model.predict(X_test)
```

```
metrics.confusion_matrix(y_true=y_test, y_pred=y_pred)
```

```
array([[1091,    6],
       [  90,  365]], dtype=int64)
```

```
print("SVM accuracy = ",metrics.accuracy_score(y_test,y_pred))
```

```
SVM accuracy =  0.9381443298969072
```

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months.

Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling> Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt #Importing the libraries

df = pd.read_csv("Churn_Modelling.csv")
```

Preprocessing.

```
df.head()
```

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age |
|---|-----------|------------|----------|-------------|-----------|--------|-----|
| 0 | 1 | 15634602 | Hargrave | 619 | France | Female | 42 |
| 1 | 2 | 15647311 | Hill | 608 | Spain | Female | 41 |
| 2 | 3 | 15619304 | Onio | 502 | France | Female | 42 |
| 3 | 4 | 15701354 | Boni | 699 | France | Female | 39 |
| 4 | 5 | 15737888 | Mitchell | 850 | Spain | Female | 43 |

| | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | \ |
|---|--------|-----------|---------------|-----------|----------------|---|
| 0 | 2 | 0.00 | 1 | 1 | 1 | |
| 1 | 1 | 83807.86 | 1 | 0 | 1 | |
| 2 | 8 | 159660.80 | 3 | 1 | 0 | |
| 3 | 1 | 0.00 | 2 | 0 | 0 | |

```
4          2  125510.82          1          1          1
```

```
    EstimatedSalary  Exited
0         101348.88        1
1         112542.58        0
2         113931.57        1
3          93826.63        0
4          79084.10        0
```

```
df.shape
```

```
(10000, 14)
```

```
df.describe()
```

```
      RowNumber  CustomerId  CreditScore  Age
Tenure \
count  10000.000000  1.000000e+04  10000.000000  10000.000000
10000.000000
mean    5000.50000  1.569094e+07    650.528800    38.921800
5.012800
std     2886.89568  7.193619e+04     96.653299    10.487806
2.892174
min         1.00000  1.556570e+07    350.000000    18.000000
0.000000
25%     2500.75000  1.562853e+07    584.000000    32.000000
3.000000
50%     5000.50000  1.569074e+07    652.000000    37.000000
5.000000
75%     7500.25000  1.575323e+07    718.000000    44.000000
7.000000
max    10000.00000  1.581569e+07    850.000000    92.000000
10.000000
```

```
      Balance  NumOfProducts  HasCrCard  IsActiveMember  \
count  10000.000000  10000.000000  10000.000000  10000.000000
mean    76485.889288    1.530200    0.70550    0.515100
std     62397.405202    0.581654    0.45584    0.499797
min         0.000000    1.000000    0.000000    0.000000
25%         0.000000    1.000000    0.000000    0.000000
50%     97198.540000    1.000000    1.000000    1.000000
75%    127644.240000    2.000000    1.000000    1.000000
max    250898.090000    4.000000    1.000000    1.000000
```

```
      EstimatedSalary  Exited
count  10000.000000  10000.000000
mean    100090.239881    0.203700
std     57510.492818    0.402769
min         11.580000    0.000000
25%     51002.110000    0.000000
50%    100193.915000    0.000000
```

| | | |
|-----|---------------|----------|
| 75% | 149388.247500 | 0.000000 |
| max | 199992.480000 | 1.000000 |

df.isnull()

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender |
|-------|-----------|------------|---------|-------------|-----------|--------|
| Age \ | | | | | | |
| 0 | False | False | False | False | False | False |
| False | | | | | | |
| 1 | False | False | False | False | False | False |
| False | | | | | | |
| 2 | False | False | False | False | False | False |
| False | | | | | | |
| 3 | False | False | False | False | False | False |
| False | | | | | | |
| 4 | False | False | False | False | False | False |
| False | | | | | | |
| ... | ... | ... | ... | ... | ... | ... |
| ... | | | | | | |
| 9995 | False | False | False | False | False | False |
| False | | | | | | |
| 9996 | False | False | False | False | False | False |
| False | | | | | | |
| 9997 | False | False | False | False | False | False |
| False | | | | | | |
| 9998 | False | False | False | False | False | False |
| False | | | | | | |
| 9999 | False | False | False | False | False | False |
| False | | | | | | |

| | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | \ |
|------|--------|---------|---------------|-----------|----------------|-------|
| 0 | False | False | False | False | False | False |
| 1 | False | False | False | False | False | False |
| 2 | False | False | False | False | False | False |
| 3 | False | False | False | False | False | False |
| 4 | False | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... | ... |
| 9995 | False | False | False | False | False | False |
| 9996 | False | False | False | False | False | False |
| 9997 | False | False | False | False | False | False |
| 9998 | False | False | False | False | False | False |
| 9999 | False | False | False | False | False | False |

| | EstimatedSalary | Exited |
|-----|-----------------|--------|
| 0 | False | False |
| 1 | False | False |
| 2 | False | False |
| 3 | False | False |
| 4 | False | False |
| ... | ... | ... |

| | | |
|------|-------|-------|
| 9995 | False | False |
| 9996 | False | False |
| 9997 | False | False |
| 9998 | False | False |
| 9999 | False | False |

[10000 rows x 14 columns]

df.isnull().sum()

| | |
|-----------------|---|
| RowNumber | 0 |
| CustomerId | 0 |
| Surname | 0 |
| CreditScore | 0 |
| Geography | 0 |
| Gender | 0 |
| Age | 0 |
| Tenure | 0 |
| Balance | 0 |
| NumOfProducts | 0 |
| HasCrCard | 0 |
| IsActiveMember | 0 |
| EstimatedSalary | 0 |
| Exited | 0 |

dtype: int64

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10000 entries, 0 to 9999

Data columns (total 14 columns):

| # | Column | Non-Null | Count | Dtype |
|----|-----------------|----------|----------|---------|
| 0 | RowNumber | 10000 | non-null | int64 |
| 1 | CustomerId | 10000 | non-null | int64 |
| 2 | Surname | 10000 | non-null | object |
| 3 | CreditScore | 10000 | non-null | int64 |
| 4 | Geography | 10000 | non-null | object |
| 5 | Gender | 10000 | non-null | object |
| 6 | Age | 10000 | non-null | int64 |
| 7 | Tenure | 10000 | non-null | int64 |
| 8 | Balance | 10000 | non-null | float64 |
| 9 | NumOfProducts | 10000 | non-null | int64 |
| 10 | HasCrCard | 10000 | non-null | int64 |
| 11 | IsActiveMember | 10000 | non-null | int64 |
| 12 | EstimatedSalary | 10000 | non-null | float64 |
| 13 | Exited | 10000 | non-null | int64 |

dtypes: float64(2), int64(9), object(3)

memory usage: 1.1+ MB

df.dtypes

```

RowNumber      int64
CustomerId     int64
Surname        object
CreditScore    int64
Geography      object
Gender         object
Age            int64
Tenure         int64
Balance        float64
NumOfProducts int64
HasCrCard      int64
IsActiveMember int64
EstimatedSalary float64
Exited         int64
dtype: object

```

```
df.columns
```

```

Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore',
      'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
      'HasCrCard',
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')

```

```

df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis= 1)
#Dropping the unnecessary columns

```

```
df.head()
```

| | CreditScore | Geography | Gender | Age | Tenure | Balance |
|---|-------------|-----------|--------|-----|--------|-----------|
| 0 | 619 | France | Female | 42 | 2 | 0.00 |
| 1 | | | | | | |
| 1 | 608 | Spain | Female | 41 | 1 | 83807.86 |
| 1 | | | | | | |
| 2 | 502 | France | Female | 42 | 8 | 159660.80 |
| 3 | | | | | | |
| 3 | 699 | France | Female | 39 | 1 | 0.00 |
| 2 | | | | | | |
| 4 | 850 | Spain | Female | 43 | 2 | 125510.82 |
| 1 | | | | | | |


| | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|---|-----------|----------------|-----------------|--------|
| 0 | 1 | 1 | 101348.88 | 1 |
| 1 | 0 | 1 | 112542.58 | 0 |
| 2 | 1 | 0 | 113931.57 | 1 |
| 3 | 0 | 0 | 93826.63 | 0 |
| 4 | 1 | 1 | 79084.10 | 0 |

Visualization

```
def visualization(x, y, xlabel):  
    plt.figure(figsize=(10,5))  
    plt.hist([x, y], color=['red', 'green'], label = ['exit',  
'not_exit'])  
    plt.xlabel(xlabel, fontsize=20)  
    plt.ylabel("No. of customers", fontsize=20)  
    plt.legend()
```


```
df_churn_exited = df[df['Exited']==1]['Tenure']  
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

```
visualization(df_churn_exited, df_churn_not_exited, "Tenure")
```



```
df_churn_exited2 = df[df['Exited']==1]['Age']  
df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

```
visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```



Converting the Categorical Variables

```
X =
df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']]
states = pd.get_dummies(df['Geography'], drop_first = True)
gender = pd.get_dummies(df['Gender'], drop_first = True)
```

```
df = pd.concat([df, gender, states], axis = 1)
```

Splitting the training and testing Dataset

```
df.head()
```

| | CreditScore | Geography | Gender | Age | Tenure | Balance |
|---|-------------|-----------|--------|-----|--------|-----------|
| 0 | 619 | France | Female | 42 | 2 | 0.00 |
| 1 | 608 | Spain | Female | 41 | 1 | 83807.86 |
| 2 | 502 | France | Female | 42 | 8 | 159660.80 |
| 3 | 699 | France | Female | 39 | 1 | 0.00 |
| 4 | 850 | Spain | Female | 43 | 2 | 125510.82 |

| | HasCrCard | IsActiveMember | EstimatedSalary | Exited | Male | Germany |
|---|-----------|----------------|-----------------|--------|------|---------|
| 0 | 1 | 1 | 101348.88 | 1 | 0 | 0 |

| | | | | | | |
|---|---|---|-----------|---|---|---|
| 1 | 0 | 1 | 112542.58 | 0 | 0 | 0 |
| 1 | | | | | | |
| 2 | 1 | 0 | 113931.57 | 1 | 0 | 0 |
| 0 | | | | | | |
| 3 | 0 | 0 | 93826.63 | 0 | 0 | 0 |
| 0 | | | | | | |
| 4 | 1 | 1 | 79084.10 | 0 | 0 | 0 |
| 1 | | | | | | |

```
X =
df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
    'IsActiveMember', 'EstimatedSalary', 'Male', 'Germany', 'Spain']]

y = df['Exited']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30)
```

Normalizing the values with mean as 0 and Standard Deviation as 1

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()

X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

X_train
array([[ 0.30685022, -0.36409498,  1.37874982, ...,  0.92216229,
        -0.5821891 ,  1.76089794],
       [ 0.47071357,  0.40957576, -0.34913058, ...,  0.92216229,
        1.71765497, -0.56789208],
       [-1.23961016, -0.36409498, -0.0035545 , ...,  0.92216229,
        -0.5821891 ,  1.76089794],
       ...,
       [ 0.45023065,  0.02274039, -0.0035545 , ...,  0.92216229,
        -0.5821891 , -0.56789208],
       [-0.0413594 , -0.84763919, -0.0035545 , ..., -1.08440782,
        -0.5821891 , -0.56789208],
       [-1.2293687 , -1.3311834 ,  1.37874982, ...,  0.92216229,
        -0.5821891 , -0.56789208]])

X_test
array([[ 0.2863673 ,  0.98982882, -1.04028274, ...,  0.92216229,
        1.71765497, -0.56789208],
       [ 0.80868173, -0.46080382,  1.37874982, ..., -1.08440782,
        -0.5821891 ,  1.76089794],
       [-0.13353254,  1.18324651, -1.38585882, ...,  0.92216229,
        1.71765497, -0.56789208],
       ...,
       ...])
```

```

[-0.31787881,  1.8602084 , -0.0035545 , ...,  0.92216229,
 -0.5821891 , -0.56789208],
[ 0.81892319,  2.24704378, -1.04028274, ..., -1.08440782,
 -0.5821891 ,  1.76089794],
[-0.51246654, -0.36409498, -0.34913058, ...,  0.92216229,
 -0.5821891 , -0.56789208]])

```

Building the Classifier Model using Keras

import keras *#Keras is the wrapper on the top of tensorflow*
#Can use Tensorflow as well but won't be able to understand the errors initially.

from keras.models import Sequential *#To create sequential neural network*

from keras.layers import Dense *#To create hidden layers*

classifier = Sequential()

#To add the layers

#Dense helps to construct the neurons

#Input Dimension means we have 11 features

Units is to create the hidden layers

#Uniform helps to distribute the weight uniformly

classifier.add(Dense(activation = "relu",input_dim = 11,units = 6, kernel_initializer = "uniform"))

classifier.add(Dense(activation = "relu",units = 6, kernel_initializer = "uniform")) *#Adding second hidden layers*

classifier.add(Dense(activation = "sigmoid",units = 1, kernel_initializer = "uniform")) *#Final neuron will be having sigmoid function*

classifier.compile(optimizer="adam",loss = 'binary_crossentropy',metrics = ['accuracy']) *#To compile the Artificial Neural Network. Used Binary crossentropy as we just have only two output*

classifier.summary() *#3 layers created. 6 neurons in 1st,6neurons in 2nd layer and 1 neuron in last*

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense (Dense) | (None, 6) | 72 |
| dense_1 (Dense) | (None, 6) | 42 |
| dense_2 (Dense) | (None, 1) | 7 |

```
=====
Total params: 121
Trainable params: 121
Non-trainable params: 0
=====
```

```
classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the
ANN to training dataset
```

```
Epoch 1/50
700/700 [=====] - 1s 652us/step - loss:
0.4923 - accuracy: 0.7966
Epoch 2/50
700/700 [=====] - 1s 763us/step - loss:
0.4267 - accuracy: 0.7966
Epoch 3/50
700/700 [=====] - 0s 622us/step - loss:
0.4225 - accuracy: 0.7966
Epoch 4/50
700/700 [=====] - 0s 651us/step - loss:
0.4182 - accuracy: 0.8129
Epoch 5/50
700/700 [=====] - 0s 645us/step - loss:
0.4156 - accuracy: 0.8253
Epoch 6/50
700/700 [=====] - 0s 655us/step - loss:
0.4132 - accuracy: 0.8294
Epoch 7/50
700/700 [=====] - 0s 643us/step - loss:
0.4117 - accuracy: 0.8316
Epoch 8/50
700/700 [=====] - 0s 633us/step - loss:
0.4102 - accuracy: 0.8331
Epoch 9/50
700/700 [=====] - 0s 649us/step - loss:
0.4093 - accuracy: 0.8324
Epoch 10/50
700/700 [=====] - 0s 662us/step - loss:
0.4077 - accuracy: 0.8346
Epoch 11/50
700/700 [=====] - 0s 680us/step - loss:
0.4071 - accuracy: 0.8343
Epoch 12/50
700/700 [=====] - 1s 787us/step - loss:
0.4064 - accuracy: 0.8356
Epoch 13/50
700/700 [=====] - 1s 765us/step - loss:
0.4054 - accuracy: 0.8353
Epoch 14/50
700/700 [=====] - 1s 754us/step - loss:
```

0.4050 - accuracy: 0.8366
Epoch 15/50
700/700 [=====] - 1s 744us/step - loss:
0.4042 - accuracy: 0.8360
Epoch 16/50
700/700 [=====] - 1s 773us/step - loss:
0.4038 - accuracy: 0.8357
Epoch 17/50
700/700 [=====] - 1s 789us/step - loss:
0.4037 - accuracy: 0.8349
Epoch 18/50
700/700 [=====] - 1s 763us/step - loss:
0.4031 - accuracy: 0.8366
Epoch 19/50
700/700 [=====] - 1s 757us/step - loss:
0.4030 - accuracy: 0.8363
Epoch 20/50
700/700 [=====] - 1s 751us/step - loss:
0.4024 - accuracy: 0.8360
Epoch 21/50
700/700 [=====] - 1s 774us/step - loss:
0.4020 - accuracy: 0.8360
Epoch 22/50
700/700 [=====] - 1s 771us/step - loss:
0.4019 - accuracy: 0.8331
Epoch 23/50
700/700 [=====] - 1s 752us/step - loss:
0.4021 - accuracy: 0.8357
Epoch 24/50
700/700 [=====] - 1s 752us/step - loss:
0.4015 - accuracy: 0.8363
Epoch 25/50
700/700 [=====] - 1s 771us/step - loss:
0.4013 - accuracy: 0.8339
Epoch 26/50
700/700 [=====] - 1s 763us/step - loss:
0.4010 - accuracy: 0.8337
Epoch 27/50
700/700 [=====] - 1s 755us/step - loss:
0.4008 - accuracy: 0.8369
Epoch 28/50
700/700 [=====] - 1s 758us/step - loss:
0.4003 - accuracy: 0.8364
Epoch 29/50
700/700 [=====] - 1s 759us/step - loss:
0.4008 - accuracy: 0.8349
Epoch 30/50
700/700 [=====] - 1s 775us/step - loss:
0.4007 - accuracy: 0.8354
Epoch 31/50

700/700 [=====] - 1s 748us/step - loss:
0.3997 - accuracy: 0.8371
Epoch 32/50
700/700 [=====] - 1s 757us/step - loss:
0.4001 - accuracy: 0.8331
Epoch 33/50
700/700 [=====] - 1s 770us/step - loss:
0.3995 - accuracy: 0.8351
Epoch 34/50
700/700 [=====] - 1s 776us/step - loss:
0.3999 - accuracy: 0.8359
Epoch 35/50
700/700 [=====] - 1s 782us/step - loss:
0.3990 - accuracy: 0.8366
Epoch 36/50
700/700 [=====] - 1s 769us/step - loss:
0.3997 - accuracy: 0.8359
Epoch 37/50
700/700 [=====] - 1s 757us/step - loss:
0.3992 - accuracy: 0.8357
Epoch 38/50
700/700 [=====] - 1s 774us/step - loss:
0.3991 - accuracy: 0.8347
Epoch 39/50
700/700 [=====] - 1s 763us/step - loss:
0.3983 - accuracy: 0.8347
Epoch 40/50
700/700 [=====] - 1s 751us/step - loss:
0.3982 - accuracy: 0.8353
Epoch 41/50
700/700 [=====] - 1s 765us/step - loss:
0.3988 - accuracy: 0.8354
Epoch 42/50
700/700 [=====] - 1s 744us/step - loss:
0.3982 - accuracy: 0.8339
Epoch 43/50
700/700 [=====] - 1s 718us/step - loss:
0.3984 - accuracy: 0.8389
Epoch 44/50
700/700 [=====] - 1s 789us/step - loss:
0.3982 - accuracy: 0.8369
Epoch 45/50
700/700 [=====] - 1s 749us/step - loss:
0.3976 - accuracy: 0.8336
Epoch 46/50
700/700 [=====] - 1s 761us/step - loss:
0.3983 - accuracy: 0.8346
Epoch 47/50
700/700 [=====] - 1s 751us/step - loss:
0.3980 - accuracy: 0.8354

```
Epoch 48/50
700/700 [=====] - 1s 757us/step - loss:
0.3980 - accuracy: 0.8353
Epoch 49/50
700/700 [=====] - 1s 764us/step - loss:
0.3981 - accuracy: 0.8349
Epoch 50/50
700/700 [=====] - 1s 746us/step - loss:
0.3979 - accuracy: 0.8357
```

```
<keras.callbacks.History at 0x2109341ca00>
```

```
y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5) #Predicting the result

from sklearn.metrics import
confusion_matrix, accuracy_score, classification_report

cm = confusion_matrix(y_test, y_pred)


cm
array([[2300,   87],
       [ 407,  206]], dtype=int64)

accuracy = accuracy_score(y_test, y_pred)

accuracy
0.8353333333333334

plt.figure(figsize = (10,7))
sns.heatmap(cm, annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')

Text(69.0, 0.5, 'Truth')
```

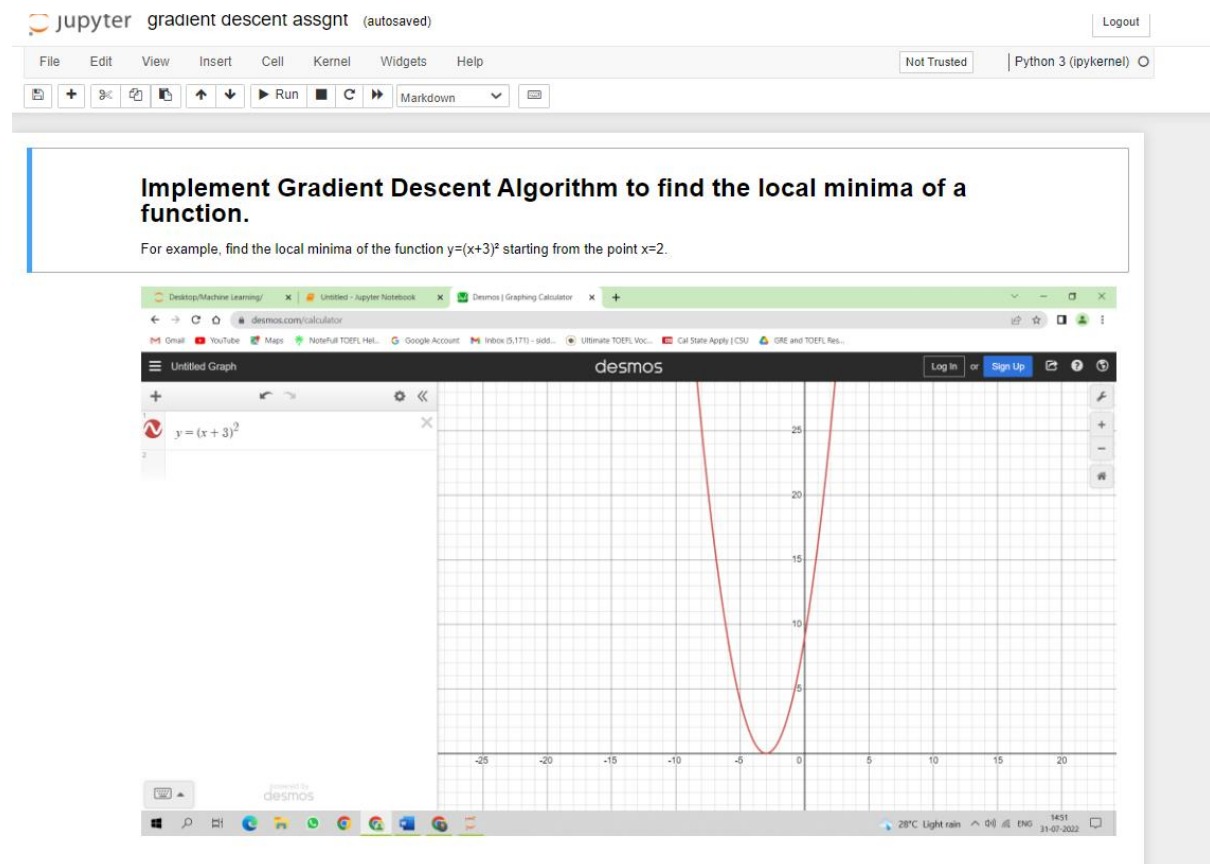


```
print(classification_report(y_test,y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.96 | 0.90 | 2387 |
| 1 | 0.70 | 0.34 | 0.45 | 613 |
| accuracy | | | 0.84 | 3000 |
| macro avg | 0.78 | 0.65 | 0.68 | 3000 |
| weighted avg | 0.82 | 0.84 | 0.81 | 3000 |

Assignment 4 :

Code:



We know the answer just by looking at the graph. $y = (x+3)^2$ reaches it's minimum value when $x = -3$ (i.e when $x=-3$, $y=0$). Hence $x=-3$ is the local and global minima of the function. Below is the implementation in python

```
In [1]: current_x = 2 # The algorithm starts at x=3
rate = 0.01 # Learning rate
precision = 0.000001 #This tells us when to stop the algorithm
previous_step_size = 1
max_iters = 10000 # maximum number of iterations
iters = 0 #iteration counter
df = lambda x: 2*(x+3) #Gradient of our function

In [2]: while previous_step_size > precision and iters < max_iters: #When Previous Step Size will be Less than Precision then we will rec
previous_x = current_x #Store current x value in prev_x
current_x = current_x - rate * df(previous_x) #Grad descent
previous_step_size = abs(current_x - previous_x) #Change in x
iters = iters+1 #iteration count
print("Iteration",iters,"\\nX value is",current_x) #Print iterations

print("The local minimum occurs at", current_x)
```

```
print("The local minimum occurs at", current_x)
```

```
Iteration 563  
X value is -2.999942555213562  
Iteration 564  
X value is -2.999943704109291  
Iteration 565  
X value is -2.999944830027105  
Iteration 566  
X value is -2.999945933426563  
Iteration 567  
X value is -2.999947014758032  
Iteration 568  
X value is -2.9999480744628713  
Iteration 569  
X value is -2.999949112973614  
Iteration 570  
X value is -2.999950130714142  
Iteration 571  
X value is -2.999951128099859  
The local minimum occurs at -2.999951128099859
```

In []:

In []:

Assignment 5

KNN algorithm on diabetes dataset

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn import metrics

df=pd.read_csv('diabetes.csv')

df.columns

Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
       'Insulin',
       'BMI', 'Pedigree', 'Age', 'Outcome'],
      dtype='object')
```

Check for null values. If present remove null values from the dataset

```
df.isnull().sum()
```

```
Pregnancies    0
Glucose         0
BloodPressure   0
SkinThickness   0
Insulin         0
BMI             0
Pedigree        0
Age            0
Outcome         0
dtype: int64
```

Outcome is the label/target, other columns are features

```
X = df.drop('Outcome',axis = 1)
y = df['Outcome']

from sklearn.preprocessing import scale
X = scale(X)
# split into train and test
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42)
```

```
from sklearn.neighbors import KNeighborsClassifier  
knn = KNeighborsClassifier(n_neighbors=7)
```

```
knn.fit(X_train, y_train)  
y_pred = knn.predict(X_test)
```

```
print("Confusion matrix: ")  
cs = metrics.confusion_matrix(y_test,y_pred)  
print(cs)
```

```
Confusion matrix:  
[[123  28]  
 [ 37  43]]
```

```
print("Accuracy ",metrics.accuracy_score(y_test,y_pred))
```

```
Accuracy  0.7186147186147186
```

Classification error rate: proportion of instances misclassified over the whole set of instances. Error rate is calculated as the total number of two incorrect predictions (FN + FP) divided by the total number of a dataset (examples in the dataset).

Also error_rate = 1- accuracy

```
total_misclassified = cs[0,1] + cs[1,0]  
print(total_misclassified)  
total_examples = cs[0,0]+cs[0,1]+cs[1,0]+cs[1,1]  
print(total_examples)  
print("Error rate",total_misclassified/total_examples)  
print("Error rate ",1-metrics.accuracy_score(y_test,y_pred))
```

```
65
```

```
231
```

```
Error rate 0.2813852813852814
```

```
Error rate  0.2813852813852814
```

```
print("Precision score",metrics.precision_score(y_test,y_pred))
```

```
Precision score 0.6056338028169014
```

```
print("Recall score ",metrics.recall_score(y_test,y_pred))
```

```
Recall score  0.5375
```

```
print("Classification report  
",metrics.classification_report(y_test,y_pred))
```

```
Classification report          precision    recall  f1-score  
support
```

| | | | | |
|--------------|------|------|------|-----|
| 0 | 0.77 | 0.81 | 0.79 | 151 |
| 1 | 0.61 | 0.54 | 0.57 | 80 |
| accuracy | | | 0.72 | 231 |
| macro avg | 0.69 | 0.68 | 0.68 | 231 |
| weighted avg | 0.71 | 0.72 | 0.71 | 231 |


```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

import warnings
warnings.filterwarnings("ignore")

titanic_data = pd.read_csv('train.csv')
titanic_test = pd.read_csv('test.csv')
titanic_data.head()

```

```

    PassengerId  Survived  Pclass  \
0              1         0       3
1              2         1       1
2              3         1       3
3              4         1       1
4              5         0       3

```

```

                                     Name    Sex  Age
SibSp  \
0                                     Braund, Mr. Owen Harris    male  22.0
1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0
1
2                                     Heikkinen, Miss. Laina  female  26.0
0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)    female  35.0
1
4                                     Allen, Mr. William Henry    male  35.0
0

```

```

    Parch    Ticket    Fare Cabin Embarked
0      0  A/5 21171    7.2500   NaN        S
1      0    PC 17599   71.2833   C85        C
2      0 STON/O2. 3101282    7.9250   NaN        S
3      0    113803   53.1000  C123        S
4      0    373450    8.0500   NaN        S

```

```
titanic_data.shape
```

```
(891, 12)
```

```
titanic_data.describe()
```

```

count    PassengerId  Survived  Pclass    Age    SibSp  \
mean      446.000000    0.383838    2.308642    29.699118    0.523008

```

| | | | | | |
|-----|------------|----------|----------|-----------|----------|
| std | 257.353842 | 0.486592 | 0.836071 | 14.526497 | 1.102743 |
| min | 1.000000 | 0.000000 | 1.000000 | 0.420000 | 0.000000 |
| 25% | 223.500000 | 0.000000 | 2.000000 | 20.125000 | 0.000000 |
| 50% | 446.000000 | 0.000000 | 3.000000 | 28.000000 | 0.000000 |
| 75% | 668.500000 | 1.000000 | 3.000000 | 38.000000 | 1.000000 |
| max | 891.000000 | 1.000000 | 3.000000 | 80.000000 | 8.000000 |

| | Parch | Fare |
|-------|------------|------------|
| count | 891.000000 | 891.000000 |
| mean | 0.381594 | 32.204208 |
| std | 0.806057 | 49.693429 |
| min | 0.000000 | 0.000000 |
| 25% | 0.000000 | 7.910400 |
| 50% | 0.000000 | 14.454200 |
| 75% | 0.000000 | 31.000000 |
| max | 6.000000 | 512.329200 |

titanic_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age             714 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   object
9   Fare            891 non-null   float64
10  Cabin           204 non-null   object
11  Embarked        889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

titanic_data.isnull().sum()

| | |
|-------------|-----|
| PassengerId | 0 |
| Survived | 0 |
| Pclass | 0 |
| Name | 0 |
| Sex | 0 |
| Age | 177 |
| SibSp | 0 |
| Parch | 0 |
| Ticket | 0 |
| Fare | 0 |

```
Cabin          687
Embarked       2
dtype: int64
```

```
titanic_data = titanic_data.drop(columns='Cabin', axis = 1)
```

```
titanic_data['Age'].fillna(titanic_data['Age'].mean(), inplace= True)
```

```
print(titanic_data['Embarked'].mode()[0])
```

S

```
titanic_data['Embarked'].fillna(titanic_data['Embarked'].mode()[0],
inplace= True)
```

```
titanic_data.isnull().sum()
```

```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age             0
SibSp           0
Parch           0
Ticket         0
Fare            0
Embarked        0
dtype: int64
```

```
titanic_data.shape
```

```
(891, 11)
```

```
titanic_data.corr()
```

| | PassengerId | Survived | Pclass | Age | SibSp |
|-------------|-------------|-----------|-----------|-----------|-----------|
| Parch \ | | | | | |
| PassengerId | 1.000000 | -0.005007 | -0.035144 | 0.033207 | -0.057527 |
| 0.001652 | | | | | |
| Survived | -0.005007 | 1.000000 | -0.338481 | -0.069809 | -0.035322 |
| 0.081629 | | | | | |
| Pclass | -0.035144 | -0.338481 | 1.000000 | -0.331339 | 0.083081 |
| 0.018443 | | | | | |
| Age | 0.033207 | -0.069809 | -0.331339 | 1.000000 | -0.232625 |
| 0.179191 | | | | | |
| SibSp | -0.057527 | -0.035322 | 0.083081 | -0.232625 | 1.000000 |
| 0.414838 | | | | | |
| Parch | -0.001652 | 0.081629 | 0.018443 | -0.179191 | 0.414838 |
| 1.000000 | | | | | |
| Fare | 0.012658 | 0.257307 | -0.549500 | 0.091566 | 0.159651 |
| 0.216225 | | | | | |

```

Fare
PassengerId  0.012658
Survived     0.257307
Pclass       -0.549500
Age          0.091566
SibSp        0.159651
Parch        0.216225
Fare         1.000000

```

```
titanic_data['Survived'].value_counts()
```

```
0    549
1    342
Name: Survived, dtype: int64
```

```
titanic_data['Sex'].value_counts()
```

```
male    577
female  314
Name: Sex, dtype: int64
```

```
titanic_data.replace({'Sex':{'male':0,'female':1}}, inplace = True)
```

```
titanic_data['Embarked'].unique()
```

```
array(['S', 'C', 'Q'], dtype=object)
```

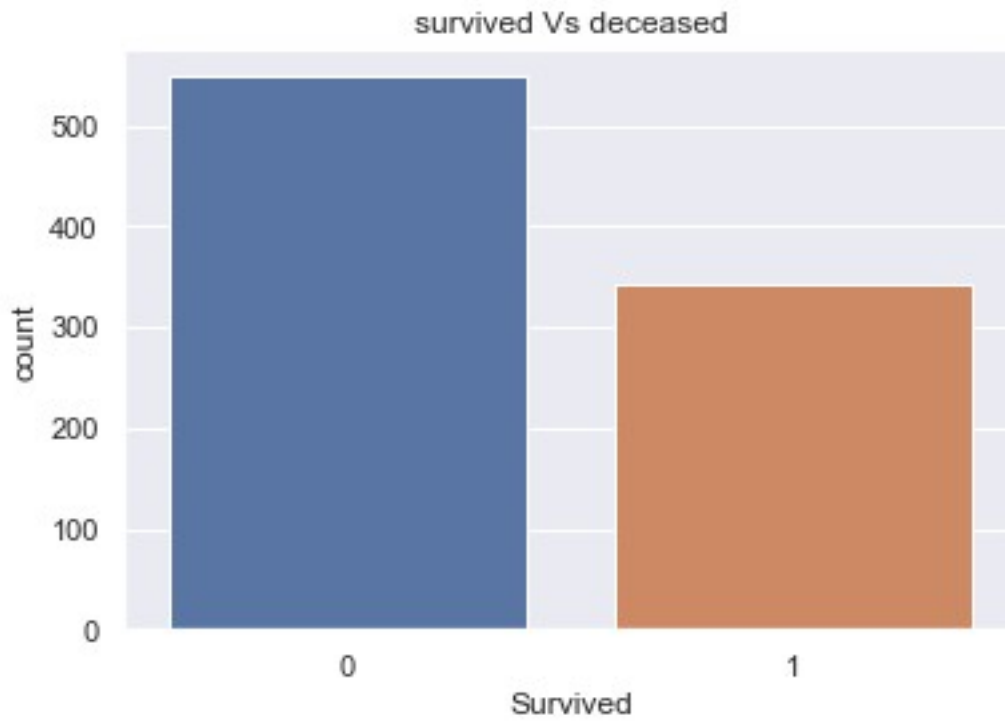
```
titanic_data.replace({'Embarked':{'S':0,'C':1, 'Q':2}}, inplace =
True)
```

```
titanic_data['Parch'].unique()
```

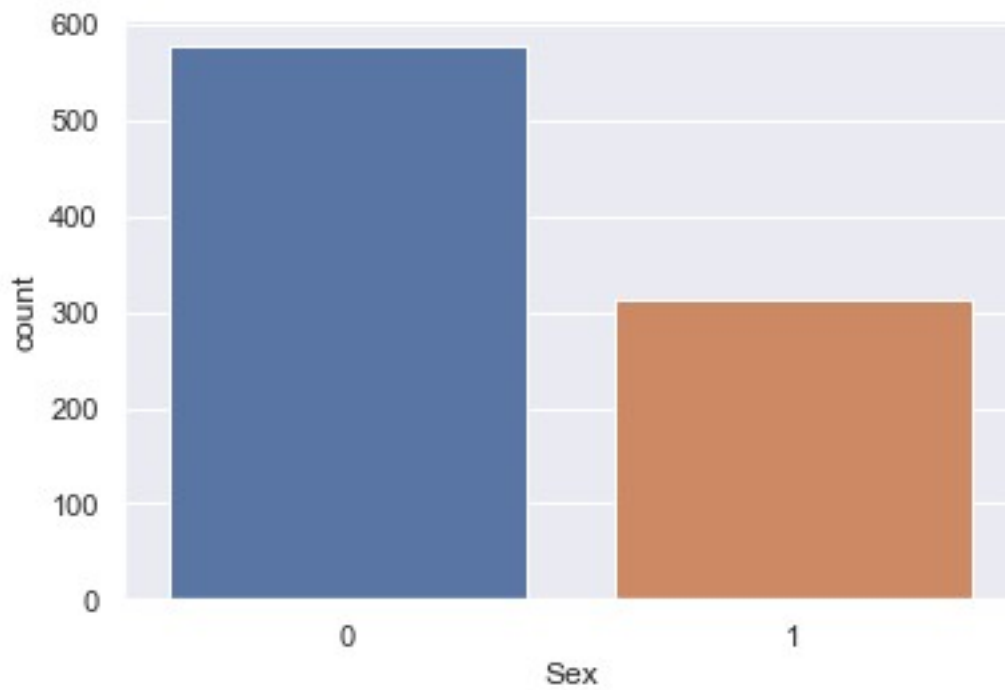
```
array([0, 1, 2, 5, 3, 4, 6], dtype=int64)
```

```
sns.set()
```

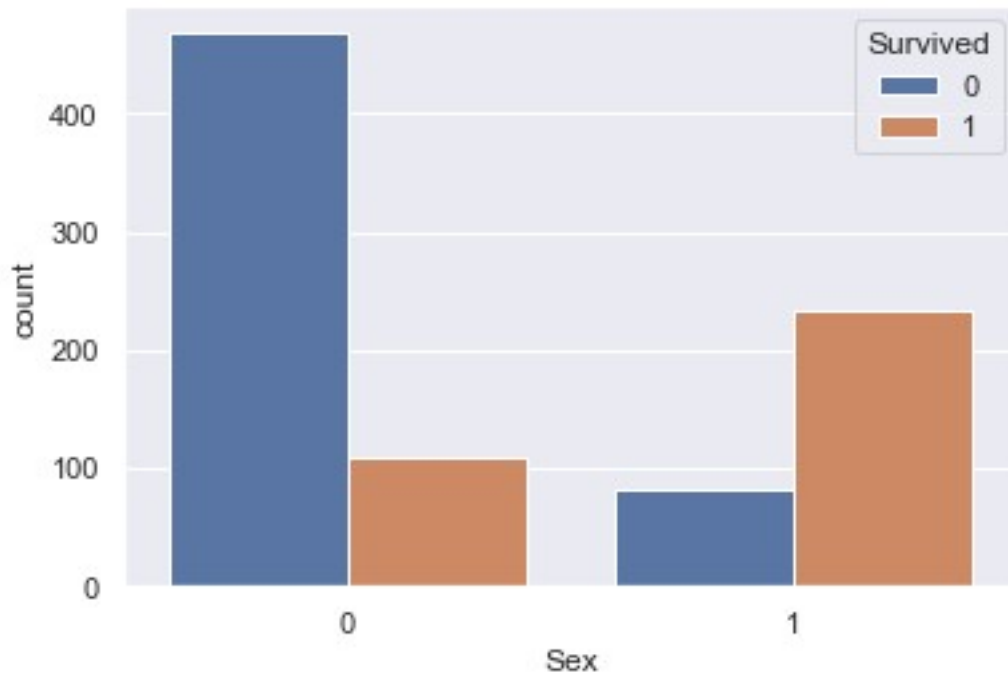
```
sns.countplot(x = titanic_data['Survived']).set_title('survived Vs
deceased');
```



```
sns.countplot(x = titanic_data['Sex']);
```



```
sns.countplot('Sex', hue='Survived', data = titanic_data);
```



```
titanic_data['Pclass'].value_counts()
```

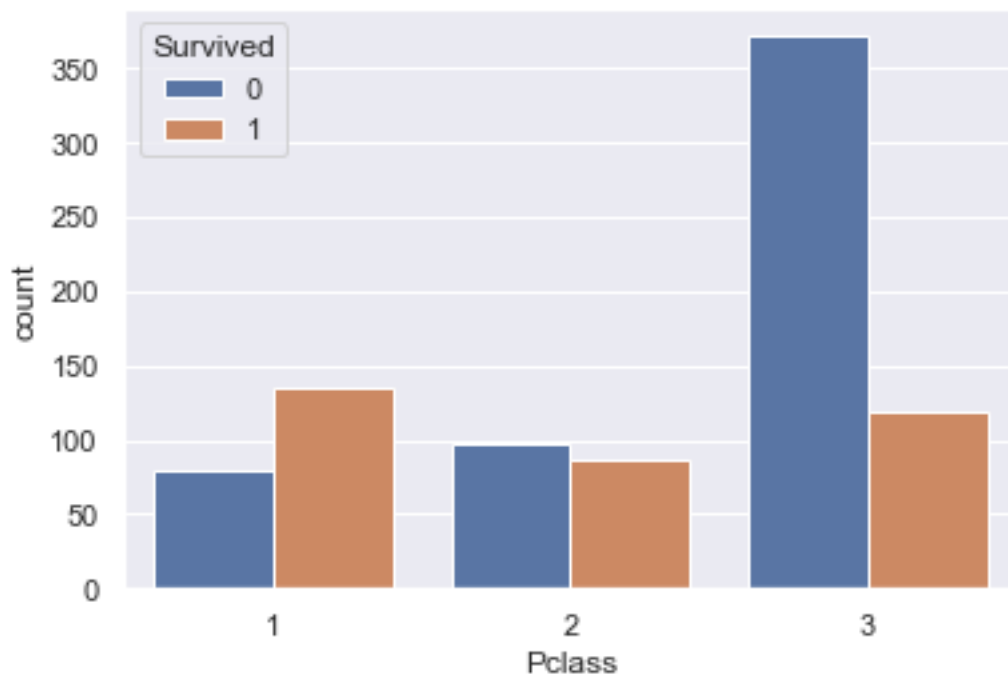
```
3    491
```

```
1    216
```

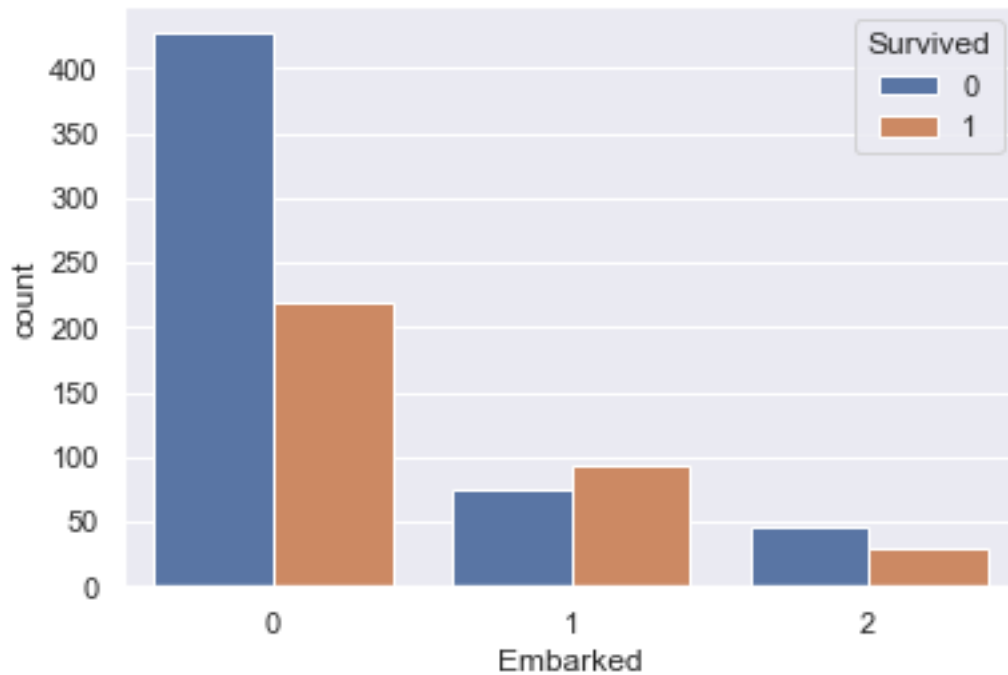
```
2    184
```

```
Name: Pclass, dtype: int64
```

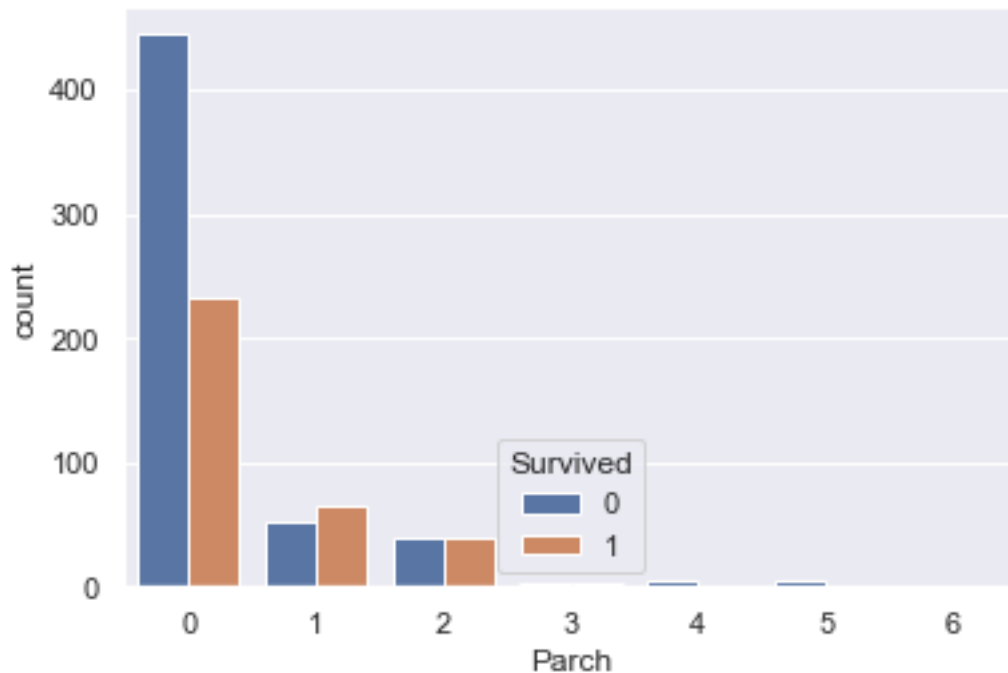
```
sns.countplot('Pclass', hue='Survived', data = titanic_data);
```



```
sns.countplot('Embarked', hue='Survived', data = titanic_data);
```



```
sns.countplot('Parch', hue='Survived', data = titanic_data);
```



```
titanic_data
```

```
0    PassengerId  Survived  Pclass  \
0         1         0         3
```

| | | | |
|-----|-----|-----|-----|
| 1 | 2 | 1 | 1 |
| 2 | 3 | 1 | 3 |
| 3 | 4 | 1 | 1 |
| 4 | 5 | 0 | 3 |
| .. | ... | ... | ... |
| 886 | 887 | 0 | 2 |
| 887 | 888 | 1 | 1 |
| 888 | 889 | 0 | 3 |
| 889 | 890 | 1 | 1 |
| 890 | 891 | 0 | 3 |

| SibSp \ | Name | Sex | Age |
|---------|---|-----|-----------|
| 0 | Braund, Mr. Owen Harris | 0 | 22.000000 |
| 1 | | | |
| 1 | Cumings, Mrs. John Bradley (Florence Briggs Th... | 1 | 38.000000 |
| 1 | | | |
| 2 | Heikkinen, Miss. Laina | 1 | 26.000000 |
| 0 | | | |
| 3 | Futrelle, Mrs. Jacques Heath (Lily May Peel) | 1 | 35.000000 |
| 1 | | | |
| 4 | Allen, Mr. William Henry | 0 | 35.000000 |
| 0 | | | |
| .. | ... | ... | ... |
| ... | | | |
| 886 | Montvila, Rev. Juozas | 0 | 27.000000 |
| 0 | | | |
| 887 | Graham, Miss. Margaret Edith | 1 | 19.000000 |
| 0 | | | |
| 888 | Johnston, Miss. Catherine Helen "Carrie" | 1 | 29.699118 |
| 1 | | | |
| 889 | Behr, Mr. Karl Howell | 0 | 26.000000 |
| 0 | | | |
| 890 | Dooley, Mr. Patrick | 0 | 32.000000 |
| 0 | | | |

| | Parch | Ticket | Fare | Embarked |
|-----|-------|------------------|---------|----------|
| 0 | 0 | A/5 21171 | 7.2500 | 0 |
| 1 | 0 | PC 17599 | 71.2833 | 1 |
| 2 | 0 | STON/O2. 3101282 | 7.9250 | 0 |
| 3 | 0 | 113803 | 53.1000 | 0 |
| 4 | 0 | 373450 | 8.0500 | 0 |
| .. | ... | ... | ... | ... |
| 886 | 0 | 211536 | 13.0000 | 0 |
| 887 | 0 | 112053 | 30.0000 | 0 |
| 888 | 2 | W./C. 6607 | 23.4500 | 0 |
| 889 | 0 | 111369 | 30.0000 | 1 |
| 890 | 0 | 370376 | 7.7500 | 2 |

[891 rows x 11 columns]


```
titanic_data.dtypes
```

```
PassengerId      int64
Survived          int64
Pclass           int64
Name             object
Sex              int64
Age             float64
SibSp            int64
Parch            int64
Ticket           object
Fare            float64
Embarked         int64
dtype: object
```

```
X = titanic_data.drop(columns=
['PassengerId', 'Name', 'Ticket', 'Survived'], axis=1)
Y = titanic_data['Survived']
```

```
print(X,Y)
```

| | Pclass | Sex | Age | SibSp | Parch | Fare | Embarked |
|-----|--------|-----|-----------|-------|-------|---------|----------|
| 0 | 3 | 0 | 22.000000 | 1 | 0 | 7.2500 | 0 |
| 1 | 1 | 1 | 38.000000 | 1 | 0 | 71.2833 | 1 |
| 2 | 3 | 1 | 26.000000 | 0 | 0 | 7.9250 | 0 |
| 3 | 1 | 1 | 35.000000 | 1 | 0 | 53.1000 | 0 |
| 4 | 3 | 0 | 35.000000 | 0 | 0 | 8.0500 | 0 |
| .. | ... | ... | ... | ... | ... | ... | ... |
| 886 | 2 | 0 | 27.000000 | 0 | 0 | 13.0000 | 0 |
| 887 | 1 | 1 | 19.000000 | 0 | 0 | 30.0000 | 0 |
| 888 | 3 | 1 | 29.699118 | 1 | 2 | 23.4500 | 0 |
| 889 | 1 | 0 | 26.000000 | 0 | 0 | 30.0000 | 1 |
| 890 | 3 | 0 | 32.000000 | 0 | 0 | 7.7500 | 2 |

```
[891 rows x 7 columns] 0      0
```

```
1      1
2      1
3      1
4      0
```

```
..
886    0
887    1
888    0
889    1
890    0
```

```
Name: Survived, Length: 891, dtype: int64
```

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=
0.2,random_state=2)
```

```
print(X_train.shape,X_test.shape,Y_train.shape,Y_test.shape)
```

```
(712, 7) (179, 7) (712,) (179,)
```

Model Training:

Logistic Regression

```
logreg = LogisticRegression()
```

```
logreg.fit(X_train,Y_train)
```

```
LogisticRegression()
```

Model Evaluation:

```
X_train_pred = logreg.predict(X_train)
```

```
X_train_pred.shape
```

```
(712,)
```

```
ac_training = accuracy_score(Y_train,X_train_pred)
```

```
print('Training Accuracy= ', round(ac_training * 100), '%')
```

```
Training Accuracy= 81 %
```

```
X_test_pred = logreg.predict(X_test)
```

```
X_test_pred.shape
```

```
(179,)
```

```
ac_testing = accuracy_score(Y_test,X_test_pred)
```

```
print('Testing Accuracy= ', round(ac_testing * 100), '%')
```

```
Testing Accuracy= 78 %
```

```
from sklearn.metrics import confusion_matrix
```

```
cf=confusion_matrix(Y_test,X_test_pred)
```

```
cf
```

```
array([[91,  9],  
       [30, 49]], dtype=int64)
```