

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months.

Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling> Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt #Importing the libraries

df = pd.read_csv("Churn_Modelling.csv")
```

Preprocessing.

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	

```
4          2  125510.82          1          1          1
```

```
    EstimatedSalary  Exited
0         101348.88        1
1         112542.58        0
2         113931.57        1
3          93826.63        0
4          79084.10        0
```

```
df.shape
```

```
(10000, 14)
```

```
df.describe()
```

```
      RowNumber  CustomerId  CreditScore  Age
Tenure \
count  10000.00000  1.000000e+04  10000.000000  10000.000000
10000.000000
mean    5000.50000  1.569094e+07    650.528800    38.921800
5.012800
std     2886.89568  7.193619e+04     96.653299    10.487806
2.892174
min         1.00000  1.556570e+07    350.000000    18.000000
0.000000
25%     2500.75000  1.562853e+07    584.000000    32.000000
3.000000
50%     5000.50000  1.569074e+07    652.000000    37.000000
5.000000
75%     7500.25000  1.575323e+07    718.000000    44.000000
7.000000
max    10000.00000  1.581569e+07    850.000000    92.000000
10.000000
```

```
      Balance  NumOfProducts  HasCrCard  IsActiveMember \
count  10000.000000  10000.000000  10000.00000  10000.000000
mean    76485.889288    1.530200    0.70550    0.515100
std     62397.405202    0.581654    0.45584    0.499797
min         0.000000    1.000000    0.00000    0.000000
25%         0.000000    1.000000    0.00000    0.000000
50%     97198.540000    1.000000    1.00000    1.000000
75%    127644.240000    2.000000    1.00000    1.000000
max    250898.090000    4.000000    1.00000    1.000000
```

```
      EstimatedSalary  Exited
count  10000.000000  10000.000000
mean    100090.239881    0.203700
std     57510.492818    0.402769
min         11.580000    0.000000
25%     51002.110000    0.000000
50%    100193.915000    0.000000
```

75%	149388.247500	0.000000
max	199992.480000	1.000000

df.isnull()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
0	False	False	False	False	False	False
False						
1	False	False	False	False	False	False
False						
2	False	False	False	False	False	False
False						
3	False	False	False	False	False	False
False						
4	False	False	False	False	False	False
False						
...
...						
9995	False	False	False	False	False	False
False						
9996	False	False	False	False	False	False
False						
9997	False	False	False	False	False	False
False						
9998	False	False	False	False	False	False
False						
9999	False	False	False	False	False	False
False						

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	False	False	False	False	False	False
1	False	False	False	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False
...
9995	False	False	False	False	False	False
9996	False	False	False	False	False	False
9997	False	False	False	False	False	False
9998	False	False	False	False	False	False
9999	False	False	False	False	False	False

	EstimatedSalary	Exited
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
...

9995	False	False
9996	False	False
9997	False	False
9998	False	False
9999	False	False

[10000 rows x 14 columns]

df.isnull().sum()

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10000 entries, 0 to 9999

Data columns (total 14 columns):

#	Column	Non-Null	Count	Dtype
0	RowNumber	10000	non-null	int64
1	CustomerId	10000	non-null	int64
2	Surname	10000	non-null	object
3	CreditScore	10000	non-null	int64
4	Geography	10000	non-null	object
5	Gender	10000	non-null	object
6	Age	10000	non-null	int64
7	Tenure	10000	non-null	int64
8	Balance	10000	non-null	float64
9	NumOfProducts	10000	non-null	int64
10	HasCrCard	10000	non-null	int64
11	IsActiveMember	10000	non-null	int64
12	EstimatedSalary	10000	non-null	float64
13	Exited	10000	non-null	int64

dtypes: float64(2), int64(9), object(3)

memory usage: 1.1+ MB

df.dtypes

```

RowNumber      int64
CustomerId      int64
Surname         object
CreditScore     int64
Geography       object
Gender          object
Age            int64
Tenure         int64
Balance        float64
NumOfProducts  int64
HasCrCard       int64
IsActiveMember int64
EstimatedSalary float64
Exited         int64
dtype: object

```

```
df.columns
```

```

Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore',
      'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
      'HasCrCard',
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')

```

```

df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis= 1)
#Dropping the unnecessary columns

```

```
df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance
0	619	France	Female	42	2	0.00
1						
1	608	Spain	Female	41	1	83807.86
1						
2	502	France	Female	42	8	159660.80
3						
3	699	France	Female	39	1	0.00
2						
4	850	Spain	Female	43	2	125510.82
1						


	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0

Visualization

```
def visualization(x, y, xlabel):  
    plt.figure(figsize=(10,5))  
    plt.hist([x, y], color=['red', 'green'], label = ['exit',  
'not_exit'])  
    plt.xlabel(xlabel, fontsize=20)  
    plt.ylabel("No. of customers", fontsize=20)  
    plt.legend()
```


```
df_churn_exited = df[df['Exited']==1]['Tenure']  
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

```
visualization(df_churn_exited, df_churn_not_exited, "Tenure")
```



```
df_churn_exited2 = df[df['Exited']==1]['Age']  
df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

```
visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```



Converting the Categorical Variables

```
X =
df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']]
states = pd.get_dummies(df['Geography'], drop_first = True)
gender = pd.get_dummies(df['Gender'], drop_first = True)
```

```
df = pd.concat([df, gender, states], axis = 1)
```

Splitting the training and testing Dataset

```
df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance
0	619	France	Female	42	2	0.00
1	608	Spain	Female	41	1	83807.86
2	502	France	Female	42	8	159660.80
3	699	France	Female	39	1	0.00
4	850	Spain	Female	43	2	125510.82

	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Male	Germany
0	1	1	101348.88	1	0	0

1	0	1	112542.58	0	0	0
1						
2	1	0	113931.57	1	0	0
0						
3	0	0	93826.63	0	0	0
0						
4	1	1	79084.10	0	0	0
1						

```
X =
df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
    'IsActiveMember', 'EstimatedSalary', 'Male', 'Germany', 'Spain']]

y = df['Exited']

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.30)
```

Normalizing the values with mean as 0 and Standard Deviation as 1

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()

X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

X_train
array([[ 0.30685022, -0.36409498,  1.37874982, ...,  0.92216229,
        -0.5821891 ,  1.76089794],
       [ 0.47071357,  0.40957576, -0.34913058, ...,  0.92216229,
        1.71765497, -0.56789208],
       [-1.23961016, -0.36409498, -0.0035545 , ...,  0.92216229,
        -0.5821891 ,  1.76089794],
       ...,
       [ 0.45023065,  0.02274039, -0.0035545 , ...,  0.92216229,
        -0.5821891 , -0.56789208],
       [-0.0413594 , -0.84763919, -0.0035545 , ..., -1.08440782,
        -0.5821891 , -0.56789208],
       [-1.2293687 , -1.3311834 ,  1.37874982, ...,  0.92216229,
        -0.5821891 , -0.56789208]])

X_test
array([[ 0.2863673 ,  0.98982882, -1.04028274, ...,  0.92216229,
        1.71765497, -0.56789208],
       [ 0.80868173, -0.46080382,  1.37874982, ..., -1.08440782,
        -0.5821891 ,  1.76089794],
       [-0.13353254,  1.18324651, -1.38585882, ...,  0.92216229,
        1.71765497, -0.56789208],
       ...,
       ...])
```



```

[-0.31787881,  1.8602084 , -0.0035545 , ...,  0.92216229,
 -0.5821891 , -0.56789208],
[ 0.81892319,  2.24704378, -1.04028274, ..., -1.08440782,
 -0.5821891 ,  1.76089794],
[-0.51246654, -0.36409498, -0.34913058, ...,  0.92216229,
 -0.5821891 , -0.56789208]])

```

Building the Classifier Model using Keras

import keras *#Keras is the wrapper on the top of tensorflow*
#Can use Tensorflow as well but won't be able to understand the errors initially.

from keras.models import Sequential *#To create sequential neural network*

from keras.layers import Dense *#To create hidden layers*

classifier = Sequential()

#To add the layers

#Dense helps to construct the neurons

#Input Dimension means we have 11 features

Units is to create the hidden layers

#Uniform helps to distribute the weight uniformly

classifier.add(Dense(activation = "relu",input_dim = 11,units = 6, kernel_initializer = "uniform"))

classifier.add(Dense(activation = "relu",units = 6, kernel_initializer = "uniform")) *#Adding second hidden layers*

classifier.add(Dense(activation = "sigmoid",units = 1, kernel_initializer = "uniform")) *#Final neuron will be having sigmoid function*

classifier.compile(optimizer="adam",loss = 'binary_crossentropy',metrics = ['accuracy']) *#To compile the Artificial Neural Network. Used Binary crossentropy as we just have only two output*

classifier.summary() *#3 layers created. 6 neurons in 1st,6neurons in 2nd layer and 1 neuron in last*

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 6)	72
dense_1 (Dense)	(None, 6)	42
dense_2 (Dense)	(None, 1)	7

```
=====
Total params: 121
Trainable params: 121
Non-trainable params: 0
```

```
classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the ANN to training dataset
```

```
Epoch 1/50
700/700 [=====] - 1s 652us/step - loss:
0.4923 - accuracy: 0.7966
Epoch 2/50
700/700 [=====] - 1s 763us/step - loss:
0.4267 - accuracy: 0.7966
Epoch 3/50
700/700 [=====] - 0s 622us/step - loss:
0.4225 - accuracy: 0.7966
Epoch 4/50
700/700 [=====] - 0s 651us/step - loss:
0.4182 - accuracy: 0.8129
Epoch 5/50
700/700 [=====] - 0s 645us/step - loss:
0.4156 - accuracy: 0.8253
Epoch 6/50
700/700 [=====] - 0s 655us/step - loss:
0.4132 - accuracy: 0.8294
Epoch 7/50
700/700 [=====] - 0s 643us/step - loss:
0.4117 - accuracy: 0.8316
Epoch 8/50
700/700 [=====] - 0s 633us/step - loss:
0.4102 - accuracy: 0.8331
Epoch 9/50
700/700 [=====] - 0s 649us/step - loss:
0.4093 - accuracy: 0.8324
Epoch 10/50
700/700 [=====] - 0s 662us/step - loss:
0.4077 - accuracy: 0.8346
Epoch 11/50
700/700 [=====] - 0s 680us/step - loss:
0.4071 - accuracy: 0.8343
Epoch 12/50
700/700 [=====] - 1s 787us/step - loss:
0.4064 - accuracy: 0.8356
Epoch 13/50
700/700 [=====] - 1s 765us/step - loss:
0.4054 - accuracy: 0.8353
Epoch 14/50
700/700 [=====] - 1s 754us/step - loss:
```

0.4050 - accuracy: 0.8366
Epoch 15/50
700/700 [=====] - 1s 744us/step - loss:
0.4042 - accuracy: 0.8360
Epoch 16/50
700/700 [=====] - 1s 773us/step - loss:
0.4038 - accuracy: 0.8357
Epoch 17/50
700/700 [=====] - 1s 789us/step - loss:
0.4037 - accuracy: 0.8349
Epoch 18/50
700/700 [=====] - 1s 763us/step - loss:
0.4031 - accuracy: 0.8366
Epoch 19/50
700/700 [=====] - 1s 757us/step - loss:
0.4030 - accuracy: 0.8363
Epoch 20/50
700/700 [=====] - 1s 751us/step - loss:
0.4024 - accuracy: 0.8360
Epoch 21/50
700/700 [=====] - 1s 774us/step - loss:
0.4020 - accuracy: 0.8360
Epoch 22/50
700/700 [=====] - 1s 771us/step - loss:
0.4019 - accuracy: 0.8331
Epoch 23/50
700/700 [=====] - 1s 752us/step - loss:
0.4021 - accuracy: 0.8357
Epoch 24/50
700/700 [=====] - 1s 752us/step - loss:
0.4015 - accuracy: 0.8363
Epoch 25/50
700/700 [=====] - 1s 771us/step - loss:
0.4013 - accuracy: 0.8339
Epoch 26/50
700/700 [=====] - 1s 763us/step - loss:
0.4010 - accuracy: 0.8337
Epoch 27/50
700/700 [=====] - 1s 755us/step - loss:
0.4008 - accuracy: 0.8369
Epoch 28/50
700/700 [=====] - 1s 758us/step - loss:
0.4003 - accuracy: 0.8364
Epoch 29/50
700/700 [=====] - 1s 759us/step - loss:
0.4008 - accuracy: 0.8349
Epoch 30/50
700/700 [=====] - 1s 775us/step - loss:
0.4007 - accuracy: 0.8354
Epoch 31/50

700/700 [=====] - 1s 748us/step - loss:
0.3997 - accuracy: 0.8371
Epoch 32/50
700/700 [=====] - 1s 757us/step - loss:
0.4001 - accuracy: 0.8331
Epoch 33/50
700/700 [=====] - 1s 770us/step - loss:
0.3995 - accuracy: 0.8351
Epoch 34/50
700/700 [=====] - 1s 776us/step - loss:
0.3999 - accuracy: 0.8359
Epoch 35/50
700/700 [=====] - 1s 782us/step - loss:
0.3990 - accuracy: 0.8366
Epoch 36/50
700/700 [=====] - 1s 769us/step - loss:
0.3997 - accuracy: 0.8359
Epoch 37/50
700/700 [=====] - 1s 757us/step - loss:
0.3992 - accuracy: 0.8357
Epoch 38/50
700/700 [=====] - 1s 774us/step - loss:
0.3991 - accuracy: 0.8347
Epoch 39/50
700/700 [=====] - 1s 763us/step - loss:
0.3983 - accuracy: 0.8347
Epoch 40/50
700/700 [=====] - 1s 751us/step - loss:
0.3982 - accuracy: 0.8353
Epoch 41/50
700/700 [=====] - 1s 765us/step - loss:
0.3988 - accuracy: 0.8354
Epoch 42/50
700/700 [=====] - 1s 744us/step - loss:
0.3982 - accuracy: 0.8339
Epoch 43/50
700/700 [=====] - 1s 718us/step - loss:
0.3984 - accuracy: 0.8389
Epoch 44/50
700/700 [=====] - 1s 789us/step - loss:
0.3982 - accuracy: 0.8369
Epoch 45/50
700/700 [=====] - 1s 749us/step - loss:
0.3976 - accuracy: 0.8336
Epoch 46/50
700/700 [=====] - 1s 761us/step - loss:
0.3983 - accuracy: 0.8346
Epoch 47/50
700/700 [=====] - 1s 751us/step - loss:
0.3980 - accuracy: 0.8354

```
Epoch 48/50
700/700 [=====] - 1s 757us/step - loss:
0.3980 - accuracy: 0.8353
Epoch 49/50
700/700 [=====] - 1s 764us/step - loss:
0.3981 - accuracy: 0.8349
Epoch 50/50
700/700 [=====] - 1s 746us/step - loss:
0.3979 - accuracy: 0.8357
```

```
<keras.callbacks.History at 0x2109341ca00>
```

```
y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5) #Predicting the result

from sklearn.metrics import
confusion_matrix, accuracy_score, classification_report

cm = confusion_matrix(y_test, y_pred)


cm
array([[2300,   87],
       [ 407,  206]], dtype=int64)

accuracy = accuracy_score(y_test, y_pred)

accuracy
0.8353333333333334

plt.figure(figsize = (10,7))
sns.heatmap(cm, annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')

Text(69.0, 0.5, 'Truth')
```



```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	2387
1	0.70	0.34	0.45	613
accuracy			0.84	3000
macro avg	0.78	0.65	0.68	3000
weighted avg	0.82	0.84	0.81	3000